Notes of [2]

Chao Tao

Feb. 28, 2020

1 Problem Setup

There is a tabular *episodic* MDP $\mathcal{M}^* = (\mathcal{S}, \mathcal{A}, \theta^*, R, H)$ where R is bounded within $[0, 1]$ and only the transition probability θ^* is *unknown*. For simplicity, we also assume the reward function R is *deterministic*. We want to find a policy such that the regret incurred by this policy after K episodes is minimized.

2 Thompson Sampling

Like *Optimism in the Face of Uncertainty*, *Thompson Sampling* dating back to [3] is another general principal guiding you how to operate in a poorly understood environment. Due to its superior empirical performance [1], it gains increasing popularity recently. Thompson Sampling is a *Bayesian* method. Basically, at the very begining, the learner equipped with this policy assumes a prior distribution $\mathcal{P}_1$ on the unknown parameter of the underlying environment i.e., θ^*. At the begining of each episode $k \geq 1$, the learner just samples a virtual environment from the posterior distribution $\mathcal{P}_k$ on θ^* which is derived based on $\mathcal{P}_{k-1}$ and the history in the $(k-1)$th episode via Bayes' Theorem and then takes the optimal policy assuming the underlying model is the sampled one. The following pseudocode shows the aforementioned learning procedure.

Algorithm 1: Thompson Sampling

```
initialization: prior distribution P_1
for episode k = 1 to K do
    compute posterior distribution P_k = P_1 | H_k
    sample θ_k from P_k and compute the optimal policy π_k
    for step h = 1 to H do
        observe state x_{k,h}
        take action a_{k,h} = π_k(x_{k,h})
```

Denote the value function starting from time t under model M' using policy π' by $V_{\pi',t}^{M'}$. Given a prior distribution $\mathcal{P}_1$ on transition probability θ^*, the *Bayesian* regret is defined by

$$\mathcal{BR}_K^{\pi} \stackrel{\text{def}}{=} \mathbb{E}_{\theta^* \sim \mathcal{P}_1}\left[\mathbb{E}\left[\sum_{k=1}^{K}(V_{*,1}^{M^*} - V_{\pi_k,1}^{M^*})(x_{k,1}) \mid \theta^*\right]\right], \tag{1}$$

where the initial state for each episode can be either randomized or *adversarial*.

3 Notations and Definitions

$[n]$	$\{1, 2, \ldots, n\}$
$\mathcal{A}$	action space
A	$\lvert\mathcal{A}\rvert$
$\mathcal{S}$	state space
S	$\lvert\mathcal{S}\rvert$
H	horizon
K	# of episodes
T	HK
$R : \mathcal{S} \times \mathcal{A} \to [0, 1]$	*known* reward function
$\theta^* : \mathcal{S} \times \mathcal{A} \to \Delta(\mathcal{S})$	transition probability of the underlying MDP
$\pi = (\pi_1, \ldots, \pi_K)$	an arbitrary policy where π_k is the policy in the kth episode
$V_{\pi',t}^{M'}$	value function starting from time t under model M' using policy π'
$x_{k,1}$	initial state of the kth episode
$(x_{k,h}, a_{k,h})$	state-action pair in the kth episode and at the hth time step
$\mathcal{H}_k$	history before the kth episode $(x_{1,1}, a_{1,1}, \ldots, x_{1,H+1}, \ldots, x_{k-1,1}, a_{k-1,1} \ldots, x_{k-1,H+1})$
$n_k(x, a)$	number of hits of state-action pair (x, a) *before* the kth episode
ρ	an arbitrary transition probability
V	an arbitrary value function
$(\rho V)(x, a)$	$\sum_{y \in \mathcal{S}} \rho(y \mid x, a) V(y)$
$\mathcal{BR}_K^{\pi}$	Bayesian regret incurred by policy π
$\mathcal{P}_k$	posterior distribution right *before* the kth episode

4 Theorem

In this lecture, we are going to show

Theorem 1. *The Bayesian regret i.e., (1) incurred by Algorithm 1 is bouned by* $\widetilde{\mathcal{O}}(HS\sqrt{AT})$.

Remark 2. *The theorem holds for any prior distribution.*

Proof. In the subsequent part, unless otherwise specified, the expectation operator is taken over all random variables. Note that θ^* is treated as a random variable. Rewrite $\mathcal{BR}_K^\pi$ we have

$$
\begin{aligned}
(1) &= \sum_{k=1}^{K} \mathbb{E}\left[(V_{*,1}^{M^*} - V_{\pi_k,1}^{M^*})(x_{k,1})\right] \\
&= \sum_{k=1}^{K} \left(\mathbb{E}\left[(V_{*,1}^{M^*} - V_{\pi_k,1}^{M_k})(x_{k,1})\right] + \mathbb{E}\left[(V_{\pi_k,1}^{M_k} - V_{\pi_k,1}^{M^*})(x_{k,1})\right]\right) \\
&= \sum_{k=1}^{K} \mathbb{E}\left[(V_{*,1}^{M^*} - V_{\pi_k,1}^{M_k})(x_{k,1})\right] + \sum_{k=1}^{K} \mathbb{E}[\widetilde{\Delta}_k], \qquad (2)
\end{aligned}
$$

where we have defined $\widetilde{\Delta}_k \stackrel{\text{def}}{=} (V_{\pi_k,1}^{M_k} - V_{\pi_k,1}^{M^*})(x_{k,1})$.

Lemma 3. $\mathbb{E}[V_{*,1}^{M^*}(x_{k,1})] = \mathbb{E}[V_{\pi_k,1}^{M_k}(x_{k,1})]$.

Proof. Just note that

$$
\begin{aligned}
\mathbb{E}[V_{*,1}^{M^*}(x_{k,1})] &= \mathbb{E}[\mathbb{E}[V_{*,1}^{M^*}(x_{k,1}) \mid \mathcal{H}_k]] \\
&= \mathbb{E}[\mathbb{E}[V_{\pi_k,1}^{M_k}(x_{k,1}) \mid \mathcal{H}_k]] \\
&= \mathbb{E}[V_{\pi_k,1}^{M_k}(x_{k,1})].
\end{aligned}
$$

□

Applying Lemma 3 in (2), we obtain

$$
(1) = \sum_{k=1}^{K} \mathbb{E}[\widetilde{\Delta}_k].
$$

Next we focus on bounding $\widetilde{\Delta}_k$. Note that

$$
\begin{aligned}
\mathbb{E}[\widetilde{\Delta}_k \mid M^*, M_k] &= \mathbb{E}[(V_{\pi_k,1}^{M_k} - V_{\pi_k,1}^{M^*})(x_{k,1}) \mid M^*, M_k] \\
&= \mathbb{E}[(\theta_k V_{\pi_k,2}^{M_k})(x_{k,1}, a_{k,1}) - (\theta^* V_{\pi_k,2}^{M^*})(x_{k,1}, a_{k,1}) \mid M^*, M_k] \\
&= \mathbb{E}[((\theta_k - \theta^*) V_{\pi_k,2}^{M_k})(x_{k,1}, a_{k,1}) + (V_{\pi_k,2}^{M_k} - V_{\pi_k,2}^{M^*})(x_{k,2}) \mid M^*, M_k] \\
&\quad + \mathbb{E}[(\theta^* (V_{\pi_k,2}^{M_k} - V_{\pi_k,2}^{M^*}))(x_{k,1}, a_{k,1}) - (V_{\pi_k,2}^{M_k} - V_{\pi_k,2}^{M^*})(x_{k,2}) \mid M^*, M_k] \\
&= \mathbb{E}[((\theta_k - \theta^*) V_{\pi_k,2}^{M_k})(x_{k,1}, a_{k,1}) + (V_{\pi_k,2}^{M_k} - V_{\pi_k,2}^{M^*})(x_{k,2}) \mid M^*, M_k],
\end{aligned}
$$

where in the second last inequality we have used

$$\mathbb{E}[(\theta^*(V_{\pi_k,2}^{M_k} - V_{\pi_k,2}^{M^*}))(x_{k,1}, a_{k,1}) - (V_{\pi_k,2}^{M_k} - V_{\pi_k,2}^{M^*})(x_{k,2}) \mid M^*, M_k] = 0.$$

By recursion and Hölder's inequality, we derive

$$\begin{aligned}
\mathbb{E}[\widetilde{\Delta}_k \mid M^*, M_k] &= \sum_{t=1}^{H} \mathbb{E}[((\theta_k - \theta^*)V_{\pi_k,t+1}^{M_k})(x_{k,t}, a_{k,t}) \mid M^*, M_k] \\
&\le \sum_{t=1}^{H} \mathbb{E}\left[\|(\theta_k - \theta^*)(x_{k,t}, a_{k,t})\|_1 \cdot \left\|V_{\pi_k,t+1}^{M_k}\right\|_\infty \mid M^*, M_k\right] \\
&\le H \cdot \sum_{t=1}^{H} \mathbb{E}[\|(\theta_k - \theta^*)(x_{k,t}, a_{k,t})\|_1 \mid M^*, M_k], \qquad (3)
\end{aligned}$$

where in the second last inequality we have used $\left\|V_{\pi_k,t+1}^{M_k}\right\|_\infty \le H$.

Let $\bar{\theta}_k(\cdot \mid s, a)$ be the empirical transition probability before the kth episode. We define $\mathcal{M}_k$ ($k \ge 2$) as the set of models such that its transition probability θ satisfies $\left\{|\bar{\theta}_k(\cdot \mid s,a) - \theta(\cdot \mid s,a)| \le C\sqrt{\frac{S\ln(SAT)}{1 \vee N_k(s,a)}}\right\}$ for all $(s,a) \in \mathcal{S} \times \mathcal{A} : N_k(s,a) > 0$ where C is a universal constant which will be defined later. According to Theorem 4, we know that there exists a constant $C > 0$ such that $\mathbf{Pr}(M_k \notin \mathcal{M}_k) \le 1/K$ and $\mathbf{Pr}(M^* \notin \mathcal{M}_k) \le 1/K$. Hence

$$\begin{aligned}
(1) &= \sum_{k=1}^{K} \mathbb{E}[\widetilde{\Delta}_k] \\
&\le \sum_{k=1}^{K} \mathbb{E}[\widetilde{\Delta}_k \mathbb{1}(M_k \in \mathcal{M}_k, M^* \in \mathcal{M}_k)] + H \cdot \sum_{k=1}^{K} (\mathbf{Pr}(M_k \notin \mathcal{M}_k) + \mathbf{Pr}(M^* \in \mathcal{M}_k)),
\end{aligned}$$

according to $\widetilde{\Delta}_k \le H$ and a union bound. Recall $\mathbf{Pr}(M_k \notin \mathcal{M}_k) \le 1/K$ and $\mathbf{Pr}(M^* \notin \mathcal{M}_k) \le 1/K$, we further obtain

$$\begin{aligned}
(1) &\lesssim \sum_{k=1}^{K} \mathbb{E}[\widetilde{\Delta}_k \mathbb{1}(M_k \in \mathcal{M}_k, M^* \in \mathcal{M}_k)] \\
&= \sum_{k=1}^{K} \mathbb{E}[\mathbb{E}[\widetilde{\Delta}_k \mathbb{1}(M_k \in \mathcal{M}_k, M^* \in \mathcal{M}_k) \mid M_k, M^*]] \\
&\le \sum_{k=1}^{K} \mathbb{E}[\mathbb{E}[\widetilde{\Delta}_k \mid M_k, M^*] \cdot \mathbb{1}(M_k \in \mathcal{M}_k, M^* \in \mathcal{M}_k)] \qquad (4)
\end{aligned}$$

Putting (3) back into (4), we have

$$
\begin{aligned}
(1) &\lesssim \sum_{k=2}^{K} \mathbb{E}\left[\sum_{t=1}^{H} \mathbb{E}[((\theta_k - \theta^*)V^{M_k}_{\pi_k,t+1})(x_{k,t}, a_{k,t}) \mid M^*, M_k] \cdot \mathbb{1}(M_k \in \mathcal{M}_k, M^* \in \mathcal{M}_k)\right] \\
&\leq \sum_{k=2}^{K} \mathbb{E}\left[\mathbb{E}\left[\sum_{h=1}^{H} CH\sqrt{\frac{S\ln(SAT)}{1 \vee N_k(x_{k,h}, a_{k,h})}} \mid M^*, M_k\right] \cdot \mathbb{1}(M_k \in \mathcal{M}_k, M^* \in \mathcal{M}_k)\right] \\
&= \mathbb{E}\left[\mathbb{E}\left[CH \cdot \sum_{k=2}^{K}\sum_{h=1}^{H} \sqrt{\frac{S\ln(SAT)}{1 \vee N_k(x_{k,h}, a_{k,h})}} \mid M^*, M_k\right] \cdot \mathbb{1}(M_k \in \mathcal{M}_k, M^* \in \mathcal{M}_k)\right], \qquad (5)
\end{aligned}
$$

where in the second last inequality we have used event $M_k \in \mathcal{M}_k$ and $M^* \in \mathcal{M}_k$.

Note that

$$
\begin{aligned}
\sum_{k=1}^{K}\sum_{h=1}^{H} \sqrt{\frac{S\ln(SAT)}{1 \vee N_k(x_{k,h}, a_{k,h})}} &\lesssim \sqrt{S\ln(SAT)} \cdot \sum_{s,a:N_K(s,a)>0} \sum_{t=1}^{N_K(s,a)} \sqrt{\frac{1}{t}} \\
&\leq 2\sqrt{S\ln(SAT)} \cdot \sum_{s,a} \sqrt{N_K(s,a)} \\
&\leq 2S\sqrt{AT\ln(SAT)} = \widetilde{\mathcal{O}}(S\sqrt{AT}) \qquad (6)
\end{aligned}
$$

Putting (6) back into (5), we prove this theorem. □

5 Tools

Theorem 4 ([4]). *Let P be a probability distribution on the set $\mathcal{S} = \{1, \ldots, S\}$. Let X_1, X_2, ..., X_m be i.i.d. random variables distributed according to P. Then, for all $\epsilon > 0$, it holds that*

$$\mathbf{Pr}(\left\|P - \bar{P}\right\|_1 \geq \epsilon) \leq (2^S - 2)\exp(-m\epsilon^2/2),$$

where $\bar{P}$ is the empirical estimation of P defined as $\bar{P}(i) = \frac{\sum_{j=1}^{m} \mathbb{1}(X_j = i)}{m}$.

References

[1] Olivier Chapelle and Lihong Li. An empirical evaluation of thompson sampling. In *NIPS*, pages 2249–2257, 2011.

[2] Ian Osband, Daniel Russo, and Benjamin Van Roy. (more) efficient reinforcement learning via posterior sampling. In *NIPS*, pages 3003–3011, 2013.

[3] William R Thompson. On the likelihood that one unknown probability exceeds another in view of the evidence of two samples. *Biometrika*, 25(3/4):285–294, 1933.

[4] Tsachy Weissman, Erik Ordentlich, Gadiel Seroussi, Sergio Verdu, and Marcelo J Weinberger. Inequalities for the l1 deviation of the empirical distribution. *Hewlett-Packard Labs, Tech. Rep*, 2003.